


The Application Layer



Slide 1

Outline

1. Network applications: services needed versus services provided
2. Example: the WEB and the HTTP
3. Example: the Internet Domain Name System (DNS)
4. Example: peer-to-peer applications (self reading)
5. Key concepts of the Socket API: demultiplexing done by UDP and TCP

1. Network applications: services needed versus services provided

1. In general, each application has
 - (a) data loss requirements
 - (b) bandwidth (throughput) requirements
 - (c) delay requirements
 - (d) security requirements
2. Data loss
 - loss-sensitive applications
 - e.g., e-mail, ftp, web documents transfers
 - for the Internet: TCP provides reliable data transfer
 - loss-tolerant applications
 - e.g., some audio/video encoding schemes generate data that can tolerate a certain amount of loss (in such cases, loss degrades the

Slide 2

Slide 3

playback quality)

- for the Internet: UDP provides fast, yet unreliable transfer mode (the media player handles transmission errors)

3. Bandwidth

■ bandwidth sensitive applications

- require a minimum amount of bandwidth to be effective
- e.g., audio: 32 Kbps to 128 Kbps (MP3 compression), video: about 1.5 Mbps (MPEG-1 compression)

■ elastic applications: make use of whatever bandwidth they get

4. Delay

■ delay sensitive applications

- e.g., Internet telephony, interactive games, etc.
- have tight constraints on the required end-to-end delay (a few hundred msec)

■ delay insensitive applications

Slide 4

5. Security

- Confidentiality: no unauthorized disclosure of information
- Integrity: data is not altered in an unauthorized way
- Authentication: sender and receiver are who/what they claim to be
- Access and availability: services are accessible and available to users

6. What services does the current Internet offer?

(a) Best effort service model; no service guarantees

(b) Connection oriented service

- provided by the TCP transport protocol
- gives a loose sense of a connection
- provides reliable, in-order, data transfer over an unreliable network
- provides flow control : sender won't overwhelm receiver
- provides congestion control : slow down sender when network is overloaded
- TCP-enhanced-with-SSL provides: confidentiality, integrity, and

Slide 5

end-point authentication

- does not provide any bandwidth or delay guarantees

(c) Connectionless service

- provided by the UDP transport protocol
- light and fast protocol: server can respond to many clients efficiently
- provides unreliable, unordered, delivery
- no flow control (so, sender can overrun receiver's buffer)
- no congestion control (so, sender can overwhelm the network)

Slide 6

2. The WEB and the HTTP


1. Section 2.2 of the textbook
2. Aside from the included multimedia content, the WEB traffic is *loss-sensitive* (hence, carried by TCP), and *elastic*.

3. Basic Components

- Clients (user agents: browsers)
- Servers (e.g., the Apache server, Sun Java System Web Server, Microsoft Internet Information Server)
- HTML pages:
 - base HTML files
 - embedded URLs of objects (other HTML files, JPEG images, Java applets, audio files, video clips, etc.)
- The HTTP (an application layer protocol)
 - defines syntax and semantics of messages, and rules of

Slide 7

communications



- currently: HTTP/1.0 (RFC 1945) and HTTP/1.1 (RFC 2616).
- HTTP is **stateless protocol** : i.e., server maintains no information about past client requests.

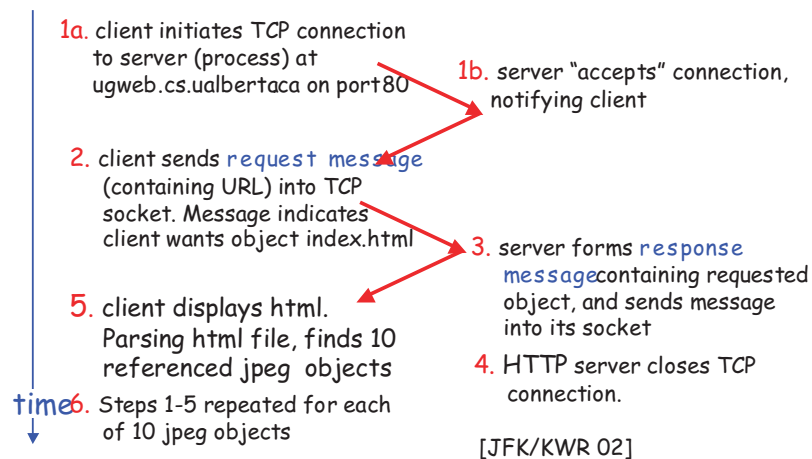
4. HTTP Connections: HTTP/1.0

- Provides **non-persistent** connections: i.e., at most one object is sent over a TCP connection. e.g., suppose user enters URL

Slide 8


ugweb.cs.ualberta.ca/index.html,

and the requested file contains references to 10 JPEG images. Then



- Response time modelling:
 - RTT (round trip time): time for a small packet travels from client to server and back

Slide 9



- total file transmission time = $2RTT + \text{transmission time}$

- [–ve] as WEB pages became more loaded with objects, browsers resorted to opening parallel TCP connections (more server overhead)

5. HTTP Connections: HTTP/1.1

- provides **persistent** connections: server leaves connection open after

sending response, subsequent HTTP messages are sent over the connection, server closes connection after timeout.

- persistent with pipelining (default mode), and persistent without pipelining

6. HTTP Messages

- Two types: **request** and **response**

- HTTP/1.0 methods:

- **GET** : client requests an object
- **POST** : client uploads data to server (possibly user data after filling a form), e.g.,

`www.somesite.com/citysearch?Canada&Edmonton`

- **HEAD** : client requests a Web page's header without the rest of its content (e.g., for indexing, or debugging, purposes)

- HTTP/1.1 additional methods:

- **PUT** : client requests uploading files in the **entity body** to a path

Slide 10

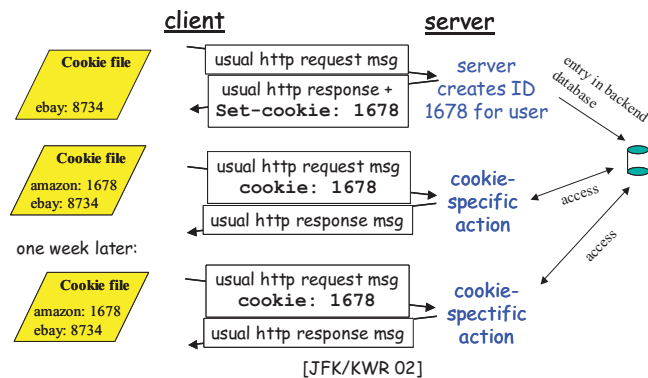
Slide 11

- specified in the URL field
- **DELETE** : client requests files to be removed from the server (e.g., for WEB publishing)
- ...

7. Keeping state using cookies

- cookies enable tracking of user activities (e.g., for maintaining “shopping carts”)
- Components:
 - cookie header line in the HTTP request/response messages
 - cookie file kept on user’s host (managed by user’s browser)
 - back-end data base at the server’s Web site

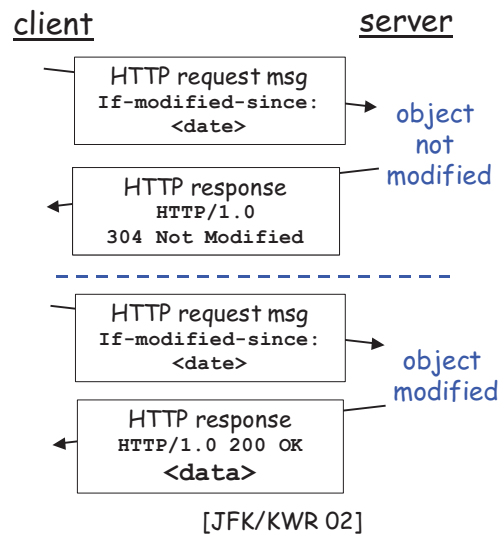
Slide 12



8. Client-side caching: conditional GET

- Goal: server does not send object, if client has an up-to-date cached version


Slide 13



9. Other Issues

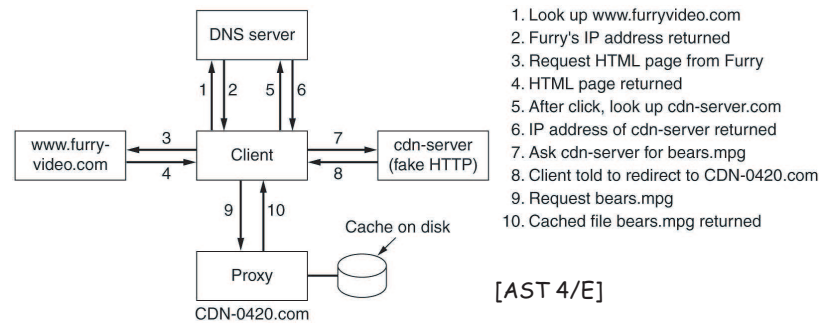
- Server-side dynamic Web page generation (e.g., CGI scripts, PHP, JavaServer Pages, ASP, etc.)

Slide 14



- Client-side dynamic Web page generation (e.g., JavaScript)
- Hierarchical caching (proxy servers)
- Server replication
- Content delivery networks

Slide 15



Slide 16

3. Example: the Internet Domain Name System (DNS)

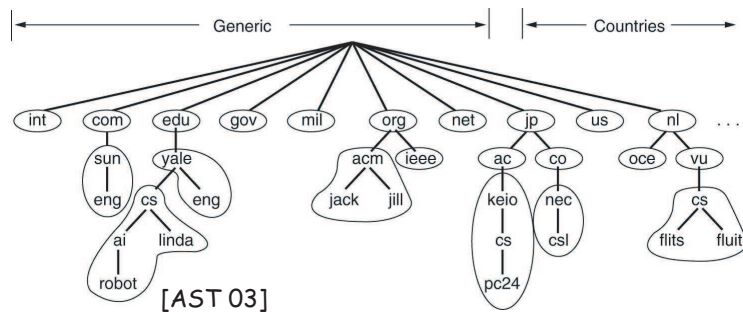
1. Section 2.5 in the textbook
2. Currently, Internet hosts and routers are identified by
 - 32-bit (IPv4) addresses (e.g. 129.128.4.241), or
 - 128-bit (IPv6) addresses
 - Human readable names: e.g., *www.ibm.com*
3. The DNS is a global distributed database system for resolving such hostname-IP address mappings. It also provides:
 - **host aliasing**: e.g., *www.ibm.com* is really *serveeast.backup2.ibm.com*
 - **mail server aliasing**: same as above (e.g., *hotmail.com* may be an alias for *relay1.west-coast.hotmial.com*)
 - **load distribution**: e.g., *cnn.com* may be replicated at 3 hosts, the DNS can rotate the host names after each reply, causing even access to the replication servers.

Slide 17

4. Note: although DNS is a core Internet service, it is implemented as a typical client-server application using the UDP transport protocol.
5. DNS is implemented using various types of **name servers** :
 - Local name servers (in close proximity of hosts)
 - A hierarchy of name servers:
 - Root name servers
 - Top-level domain (TLD) name servers
 - Authoritative name servers
6. Functionality:
 - **Local name servers**
 - organizations run one (or more) name servers (e.g., using BIND, a public-domain name server for UNIX machines)
 - application queries local servers first (e.g., an application may call `'gethostbyname () '`)
 - **Authoritative name servers**

Slide 18

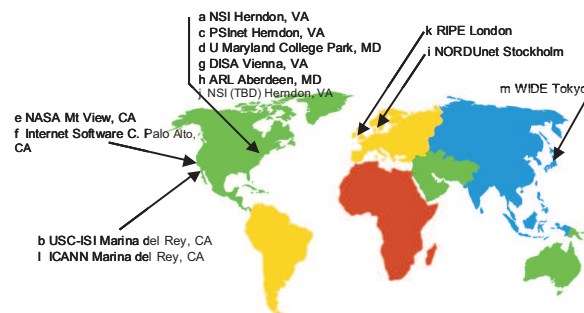
- By definition, a name server is authoritative for a host if it always have a DNS record of it.
- Each host is registered with at least two authoritative name servers (one is in the host's local ISP).
- **Top-level domain (TLD) servers**
 - Responsible for TLDs: e.g. .com, .org, .edu, .ca, .uk, .fr, etc.
 - Example companies involved at this layer:
 - * Network Solutions: servers for **.com** TLD
 - * Educause: servers for **.edu** TLD
 - For each domain (e.g., "ualberta.ca"), a TLD server stores either
 - * authoritative servers for the domain, or
 - * intermediate servers that know about authoritative servers for the domain



Slide 19

■ Root name servers

- contacted by TLD servers (and other name servers) that can not resolve a name
- about a dozen root name servers exist worldwide



[JFK/KWR 04]

13 root name
servers
worldwide

Slide 20

7. DNS Resource Records (RR)

Typical RR Fields: (name, time_to_live, class, type, value) , e.g.,

flits.cs.vu.nl.	86400	IN	HINFO	Sun Unix
flits.cs.vu.nl.	86400	IN	A	130.37.16.112
flits.cs.vu.nl.	86400	IN	A	192.31.231.165
flits.cs.vu.nl.	86400	IN	MX	1 flits.cs.vu.nl.
flits.cs.vu.nl.	86400	IN	MX	2 zephyr.cs.vu.nl.
flits.cs.vu.nl.	86400	IN	MX	3 top.cs.vu.nl.
www.cs.vu.nl.	86400	IN	CNAME	star.cs.vu.nl
ftp.cs.vu.nl.	86400	IN	CNAME	zephyr.cs.vu.nl

Slide 21

where

- **name:** is a hostname (simple such as 'flits', or fully qualified), or a domain name (e.g., 'foo.com')
- **time_to_live:** indicator of how stable the record is (highly volatile records are assigned small values)
- **class:** always 'IN' for Internet information
- **type** and **value** (some examples):
 - if (type= HINFO) value is CPU and OS information


- if (type= A) value is the IP address of the hostname
- if (type= NS) value is IP address of authoritative name server for the domain
- if (type= CNAME) value is a real name for the alias name in the first field
- if (type= MX) value is name of a mailserver associated with the hostname (first field) plus a preference number (starting with the smallest value)

Slide 22

8. Example: iterated DNS query

- in iterated query, server replies with name of yet another server to contact

Slide 23




[JFK/KWR 04]

9. **Example:** recursive DNS query

Slide 24

- recursive query puts burden on contacted name server




[JFK/KWR 04]

Slide 25

4. Key Concepts of the Socket API

1. Background

- First introduced in Unix BSD (4.2) to support the ARPA network (nowadays, an industry standard).
- Sockets as gateways to the transport layer



Slide 26

- Unix Network Programming [Stevens, Fenner, and Rudoff 2004]:



[SFR 04]

- structured around 'short' programs, each program illustrates some interesting concept
- code is available free on-line
- code uses a library that simplifies matters, without hiding details, e.g.

```
#define SA struct sockaddr

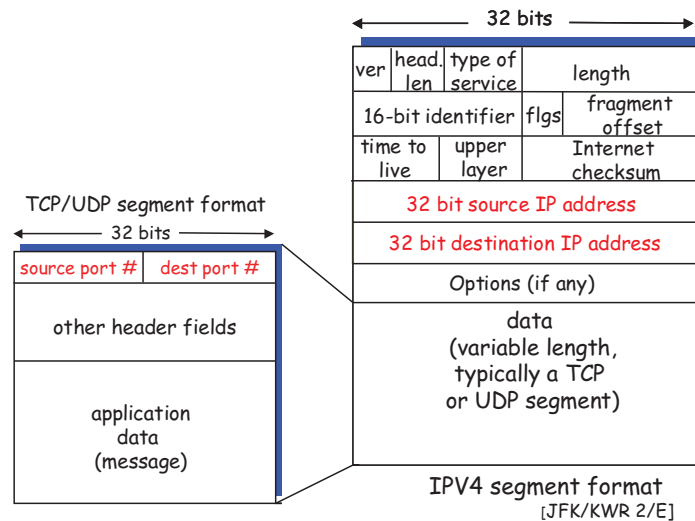
void
Connect(int fd, const struct sockaddr *sa, socklen_t salen)
{
    if (connect(fd, sa, salen) < 0)
        err_sys("connect error");
}
```

2. Demultiplexing: a transport layer service

- Among other services, the transport layer provides a logical communication channel between processes (process-to-process service)
- Using sockets, how does the transport layer find the **right process** to deliver a TCP or UDP datagram to?
 1. To start, let's look at parts of the TCP/UDP and IPv4 segment structures:

Slide 27

Slide 28



2. The client (or server) programs opens a socket using a sequence like

```
int    sockfd;
sockfd= socket (AF_INET, SOCK_STREAM, 0); // for TCP socket
sockfd= socket (AF_INET, SOCK_DGRAM, 0);  // for UDP socket
sockfd= socket (AF_INET, SOCK_RAW, 0);    // ...
```

Slide 29

and then **binds** the socket to a **port** number.

3. port numbers assigned by the **Internet Assigned Numbers Authority (IANA)**:


- **well-known ports** : 0 through 1023
- **registered ports** : 1024 through 49151 (not controlled by IANA, but kept track of by the IANA as a convenience to the community)
- **dynamic or private ports** : 49152 to 65535 (called **ephemeral** ports in [SFR 04])

4. So, **demultiplexing** at rcv host means delivering received segments to correct socket

5. **Demultiplexing done by UDP (connectionless)** :

- each UDP socket is identified by two-tuple:
(dest IP address, dest port number)

Slide 30



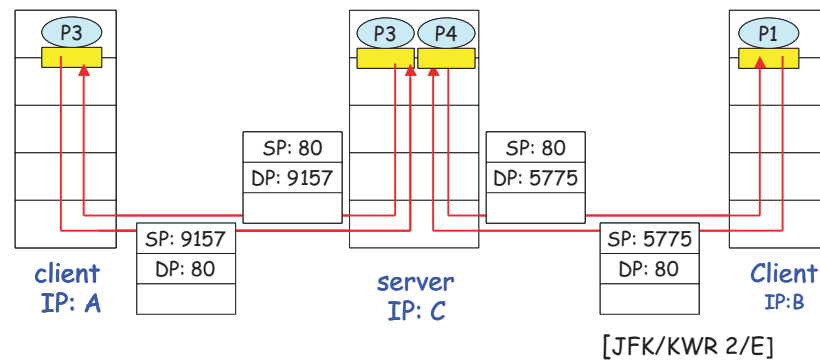
SP provides "return address"

[JFK/KWR 2/E]

6. **Demultiplexing done by TCP (connection-oriented)** :

- each TCP socket is identified by 4-tuple:
(source IP address, source port number, dest IP address, dest port number)

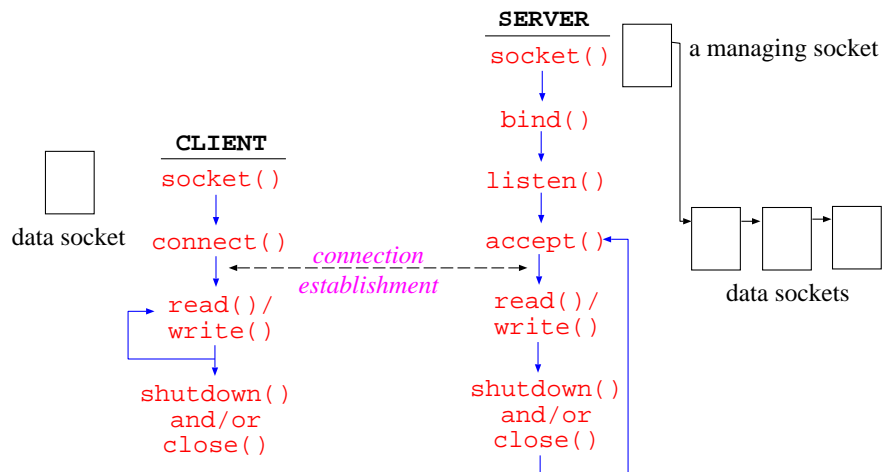
Slide 31



3. Connection Setup for TCP Sockets

- The general sequence of calls required to implement such reliable stream delivery is *asymmetric*, as shown below:

Slide 32



- In more detail:

Slide 33

